


```
In [1]:
# This Python 3 environment comes with many helpful analytics
# libraries installed
# It is defined by the kaggle/python docker image: https://git
# hub.com/kaggle/docker-python
# For example, here's several helpful packages to load in

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.r
ead_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift
+Enter) will list the files in the input directory

import os
print(os.listdir("../input"))
import keras
print(keras.__version__)
# Any results you write to the current directory are saved as
output.
```

['humpback-whale-identification', 'whales-cropped']

Using TensorFlow backend.

2.2.4

```
In [2]:
HW = 'humpback-whale-identification'
# TRAIN = '../input/humpback-whale-identification/train/'
TRAIN_CROPPED = "whales-cropped/cropped_train/cropped_train/"
TRAIN_CROPPED_IN = '../input/' + TRAIN_CROPPED

# TEST = '../input/humpback-whale-identification/test/'
TEST_CROPPED = "whales-cropped/cropped_test/cropped_test/"
TEST_CROPPED_IN = '../input/' + TEST_CROPPED

LABELS = '../input/humpback-whale-identification/train.csv'
SAMPLE_SUB = '../input/humpback-whale-identification/sample_s
ubmission.csv'

train = pd.read_csv(LABELS)
print("With new_whale:")
train.head()
```

With new_whale:

Out[2]:

	Image	Id
0	0000e88ab.jpg	w_f48451c
1	0001f9222.jpg	w_c3d896a
2	00029d126.jpg	w_20df2c5
3	00050a15a.jpg	new_whale
4	0005c1ef8.jpg	new_whale

In [3]:

```
MODEL_F = 'Model_ResNet50.h5'
WEIGHTS_F = 'Weights_ResNet50.h5'
MODEL = '../input/ResNet50/' + MODEL_F
```

```
MODEL = '../input/ResNet50/' + MODEL_F  
WEIGHTS = '../input/ResNet50/' + WEIGHTS_F
```

```
In [4]:  
    train.describe()
```

Out[4]:

	Image	Id
count	25361	25361
unique	25361	5005
top	ca07480bd.jpg	new_whale
freq	1	9664

```
In [5]:  
    import random  
    from IPython.display import Image  
    print("Example whale image")  
  
    #show sample image  
    name = random.choice(train['Image'])  
    print(name)  
    Image(filename = TRAIN_CROPPED_IN + name)
```

Example whale image
804331825.jpg

Out[5]:



```
In [6]:  
    criteria = train['Id'] != 'new_whale'  
    whales_train = train[criteria]  
  
    print("Without new_whale:")  
    whales_train.head()
```

Without new_whale:

Out[6]:

	Image	Id
0	0000e88ab.jpg	w_f48451c
1	0001f9222.jpg	w_c3d896a
2	00029d126.jpg	w_20df2c5
6	000a6daec.jpg	w_dd88965
8	0016b897a.jpg	w_64404ac

```
In [7]:  
    unique_labels = np.unique(whales_train.Id.values)  
    labels_list = unique_labels
```

```
In [8]:  
    whales_train.describe()
```

Out[8]:

	Image	Id
count	15697	15697
unique	15697	5004
top	6a28f4d47.jpg	w_23a388d
freq	1	73

In [9]:

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from matplotlib.pyplot import imshow

from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder

from keras import layers
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator

# from keras.applications.imagenet_utils import preprocess_input
from keras.applications.resnet50 import ResNet50, preprocess_input
# from keras.applications.xception import Xception, preprocess_input
# from keras.applications.inception_resnet_v2 import InceptionResNetV2, preprocess_input

from keras.losses import binary_crossentropy
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Input, Dense, Activation, BatchNormalization, Flatten, Conv2D, GlobalAveragePooling2D
from keras.layers import AveragePooling2D, MaxPooling2D, Dropout
from keras.models import Model
from keras.metrics import top_k_categorical_accuracy

import keras.backend as K
from keras.models import Sequential
from PIL import Image
import gc
import warnings
warnings.simplefilter("ignore", category=DeprecationWarning)

%matplotlib inline
```

In [10]:

```
IMAGE_HEIGHT = 128
IMAGE_WIDTH = 128
IMAGE_SHAPE = (IMAGE_HEIGHT, IMAGE_WIDTH, 3)
```

In [11]:

```
CLASSES = 5004
EPOCHS = 5
BATCH_SIZE = 32

def top_5_accuracy(y_true, y_pred):
    return top_k_categorical_accuracy(y_true, y_pred, k=5)

# setup model
base_model = ResNet50(weights='imagenet', include_top=False,
```

```

        input_shape = IMAGE_SHAPE)

x = base_model.output
x = GlobalAveragePooling2D(name='avg_pool')(x)
x = Dropout(0.4)(x)
predictions = Dense(CLASSES, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=predictions)

# transfer learning
for layer in base_model.layers:
    layer.trainable = True

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy', 'mae', top_5_accuracy])

model.summary()

```

```

/opt/conda/lib/python3.6/site-packages/keras_applications/res
net50.py:265: UserWarning: The output shape of `ResNet50(include_
ude_top=False)` has been changed since Keras 2.2.0.
    warnings.warn('The output shape of `ResNet50(include_top=False)` '

```

```

Downloading data from https://github.com/fchollet/deep-learnin
ng-models/releases/download/v0.2/resnet50_weights_tf_dim_orde
ring_tf_kernels_notop.h5
94658560/94653016 [=====] - 7s 0us/s
step
-----
```

Layer (type)	Output Shape	Param #
Connected to		
input_1 (InputLayer)	(None, 128, 128, 3)	0

```

-----
```

conv1_pad (ZeroPadding2D)	(None, 134, 134, 3)	0
input_1[0][0]		

```

-----
```

conv1 (Conv2D)	(None, 64, 64, 64)	9472
conv1_pad[0][0]		

```

-----
```

bn_conv1 (BatchNormalization)	(None, 64, 64, 64)	256
conv1[0][0]		

```

-----
```

activation_1 (Activation)	(None, 64, 64, 64)	0
bn_conv1[0][0]		

```

-----
```

pool1_pad (ZeroPadding2D)	(None, 66, 66, 64)	0
activation_1[0][0]		

```

-----
```

max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 64)	0
pool1_pad[0][0]		

```

-----
```

res2a_branch2a (Conv2D)	(None, 32, 32, 64)	4160
max_pooling2d_1[0][0]		

```
-----  
bn2a_branch2a (BatchNormalizati (None, 32, 32, 64) 256  
    res2a_branch2a[0][0]  
-----  
activation_2 (Activation)      (None, 32, 32, 64)  0  
    bn2a_branch2a[0][0]  
-----  
res2a_branch2b (Conv2D)       (None, 32, 32, 64) 36928  
    activation_2[0][0]  
-----  
bn2a_branch2b (BatchNormalizati (None, 32, 32, 64) 256  
    res2a_branch2b[0][0]  
-----  
activation_3 (Activation)      (None, 32, 32, 64)  0  
    bn2a_branch2b[0][0]  
-----  
res2a_branch2c (Conv2D)       (None, 32, 32, 256) 16640  
    activation_3[0][0]  
-----  
res2a_branch1 (Conv2D)        (None, 32, 32, 256) 16640  
    max_pooling2d_1[0][0]  
-----  
bn2a_branch2c (BatchNormalizati (None, 32, 32, 256) 1024  
    res2a_branch2c[0][0]  
-----  
bn2a_branch1 (BatchNormalizatio (None, 32, 32, 256) 1024  
    res2a_branch1[0][0]  
-----  
add_1 (Add)                  (None, 32, 32, 256)  0  
    bn2a_branch2c[0][0]  
    bn2a_branch1[0][0]  
-----  
activation_4 (Activation)      (None, 32, 32, 256)  0  
    add_1[0][0]  
-----  
res2b_branch2a (Conv2D)       (None, 32, 32, 64) 16448  
    activation_4[0][0]  
-----  
bn2b_branch2a (BatchNormalizati (None, 32, 32, 64) 256  
    res2b_branch2a[0][0]  
-----  
activation_5 (Activation)      (None, 32, 32, 64)  0  
    bn2b_branch2a[0][0]  
-----  
res2b_branch2b (Conv2D)       (None, 32, 32, 64) 36928  
    activation_5[0][0]  
-----  
bn2b_branch2b (BatchNormalizati (None, 32, 32, 64) 256  
    res2b_branch2b[0][0]
```

```
activation_6 (Activation)      (None, 32, 32, 64)  0
bn2b_branch2b[0][0]

res2b_branch2c (Conv2D)       (None, 32, 32, 256) 16640
activation_6[0][0]

bn2b_branch2c (BatchNormalizati (None, 32, 32, 256) 1024
res2b_branch2c[0][0]

add_2 (Add)                  (None, 32, 32, 256)  0
bn2b_branch2c[0][0]

activation_4[0][0]

activation_7 (Activation)      (None, 32, 32, 256)  0
add_2[0][0]

res2c_branch2a (Conv2D)       (None, 32, 32, 64)   16448
activation_7[0][0]

bn2c_branch2a (BatchNormalizati (None, 32, 32, 64)  256
res2c_branch2a[0][0]

activation_8 (Activation)      (None, 32, 32, 64)  0
bn2c_branch2a[0][0]

res2c_branch2b (Conv2D)       (None, 32, 32, 64)   36928
activation_8[0][0]

bn2c_branch2b (BatchNormalizati (None, 32, 32, 64)  256
res2c_branch2b[0][0]

activation_9 (Activation)      (None, 32, 32, 64)  0
bn2c_branch2b[0][0]

res2c_branch2c (Conv2D)       (None, 32, 32, 256)  16640
activation_9[0][0]

bn2c_branch2c (BatchNormalizati (None, 32, 32, 256) 1024
res2c_branch2c[0][0]

add_3 (Add)                  (None, 32, 32, 256)  0
bn2c_branch2c[0][0]

activation_7[0][0]

activation_10 (Activation)     (None, 32, 32, 256)  0
add_3[0][0]

res3a_branch2a (Conv2D)       (None, 16, 16, 128)  32896
```

activation_10[0][0]

bn3a_branch2a (BatchNormalizati (None, 16, 16, 128) 512
res3a_branch2a[0][0]

activation_11 (Activation) (None, 16, 16, 128) 0
bn3a_branch2a[0][0]

res3a_branch2b (Conv2D) (None, 16, 16, 128) 147584
activation_11[0][0]

bn3a_branch2b (BatchNormalizati (None, 16, 16, 128) 512
res3a_branch2b[0][0]

activation_12 (Activation) (None, 16, 16, 128) 0
bn3a_branch2b[0][0]

res3a_branch2c (Conv2D) (None, 16, 16, 512) 66048
activation_12[0][0]

res3a_branch1 (Conv2D) (None, 16, 16, 512) 131584
activation_10[0][0]

bn3a_branch2c (BatchNormalizati (None, 16, 16, 512) 2048
res3a_branch2c[0][0]

bn3a_branch1 (BatchNormalizatio (None, 16, 16, 512) 2048
res3a_branch1[0][0]

add_4 (Add) (None, 16, 16, 512) 0
bn3a_branch2c[0][0]

bn3a_branch1[0][0]

activation_13 (Activation) (None, 16, 16, 512) 0
add_4[0][0]

res3b_branch2a (Conv2D) (None, 16, 16, 128) 65664
activation_13[0][0]

bn3b_branch2a (BatchNormalizati (None, 16, 16, 128) 512
res3b_branch2a[0][0]

activation_14 (Activation) (None, 16, 16, 128) 0
bn3b_branch2a[0][0]

res3b_branch2b (Conv2D) (None, 16, 16, 128) 147584
activation_14[0][0]

bn3b_branch2b (BatchNormalizati (None, 16, 16, 128) 512

```
res3b_branch2b[0][0]

-----
activation_15 (Activation)      (None, 16, 16, 128)  0
    bn3b_branch2b[0][0]

-----
res3b_branch2c (Conv2D)        (None, 16, 16, 512)  66048
    activation_15[0][0]

-----
bn3b_branch2c (BatchNormalizati (None, 16, 16, 512)  2048
    res3b_branch2c[0][0]

-----
add_5 (Add)                  (None, 16, 16, 512)  0
    bn3b_branch2c[0][0]

    activation_13[0][0]

-----
activation_16 (Activation)      (None, 16, 16, 512)  0
    add_5[0][0]

-----
res3c_branch2a (Conv2D)        (None, 16, 16, 128)  65664
    activation_16[0][0]

-----
bn3c_branch2a (BatchNormalizati (None, 16, 16, 128)  512
    res3c_branch2a[0][0]

-----
activation_17 (Activation)      (None, 16, 16, 128)  0
    bn3c_branch2a[0][0]

-----
res3c_branch2b (Conv2D)        (None, 16, 16, 128)  147584
    activation_17[0][0]

-----
bn3c_branch2b (BatchNormalizati (None, 16, 16, 128)  512
    res3c_branch2b[0][0]

-----
activation_18 (Activation)      (None, 16, 16, 128)  0
    bn3c_branch2b[0][0]

-----
res3c_branch2c (Conv2D)        (None, 16, 16, 512)  66048
    activation_18[0][0]

-----
bn3c_branch2c (BatchNormalizati (None, 16, 16, 512)  2048
    res3c_branch2c[0][0]

-----
add_6 (Add)                  (None, 16, 16, 512)  0
    bn3c_branch2c[0][0]

    activation_16[0][0]

-----
activation_19 (Activation)      (None, 16, 16, 512)  0
    add_6[0][0]
```

res3d_branch2a (Conv2D) (None, 16, 16, 128) 65664
activation_19[0][0]

bn3d_branch2a (BatchNormalizati (None, 16, 16, 128) 512
res3d_branch2a[0][0]

activation_20 (Activation) (None, 16, 16, 128) 0
bn3d_branch2a[0][0]

res3d_branch2b (Conv2D) (None, 16, 16, 128) 147584
activation_20[0][0]

bn3d_branch2b (BatchNormalizati (None, 16, 16, 128) 512
res3d_branch2b[0][0]

activation_21 (Activation) (None, 16, 16, 128) 0
bn3d_branch2b[0][0]

res3d_branch2c (Conv2D) (None, 16, 16, 512) 66048
activation_21[0][0]

bn3d_branch2c (BatchNormalizati (None, 16, 16, 512) 2048
res3d_branch2c[0][0]

add_7 (Add) (None, 16, 16, 512) 0
bn3d_branch2c[0][0]

activation_19[0][0]

activation_22 (Activation) (None, 16, 16, 512) 0
add_7[0][0]

res4a_branch2a (Conv2D) (None, 8, 8, 256) 131328
activation_22[0][0]

bn4a_branch2a (BatchNormalizati (None, 8, 8, 256) 1024
res4a_branch2a[0][0]

activation_23 (Activation) (None, 8, 8, 256) 0
bn4a_branch2a[0][0]

res4a_branch2b (Conv2D) (None, 8, 8, 256) 590080
activation_23[0][0]

bn4a_branch2b (BatchNormalizati (None, 8, 8, 256) 1024
res4a_branch2b[0][0]

activation_24 (Activation) (None, 8, 8, 256) 0
bn4a_branch2b[0][0]

res4a_branch2c (Conv2D)	(None, 8, 8, 1024)	263168
activation_24[0][0]		
res4a_branch1 (Conv2D)	(None, 8, 8, 1024)	525312
activation_22[0][0]		
bn4a_branch2c (BatchNormalizati	(None, 8, 8, 1024)	4096
res4a_branch2c[0][0]		
bn4a_branch1 (BatchNormalizatio	(None, 8, 8, 1024)	4096
res4a_branch1[0][0]		
add_8 (Add)	(None, 8, 8, 1024)	0
bn4a_branch2c[0][0]		
bn4a_branch1[0][0]		
activation_25 (Activation)	(None, 8, 8, 1024)	0
add_8[0][0]		
res4b_branch2a (Conv2D)	(None, 8, 8, 256)	262400
activation_25[0][0]		
bn4b_branch2a (BatchNormalizati	(None, 8, 8, 256)	1024
res4b_branch2a[0][0]		
activation_26 (Activation)	(None, 8, 8, 256)	0
bn4b_branch2a[0][0]		
res4b_branch2b (Conv2D)	(None, 8, 8, 256)	590080
activation_26[0][0]		
bn4b_branch2b (BatchNormalizati	(None, 8, 8, 256)	1024
res4b_branch2b[0][0]		
activation_27 (Activation)	(None, 8, 8, 256)	0
bn4b_branch2b[0][0]		
res4b_branch2c (Conv2D)	(None, 8, 8, 1024)	263168
activation_27[0][0]		
bn4b_branch2c (BatchNormalizati	(None, 8, 8, 1024)	4096
res4b_branch2c[0][0]		
add_9 (Add)	(None, 8, 8, 1024)	0
bn4b_branch2c[0][0]		
activation_25[0][0]		
activation_28 (Activation)	(None, 8, 8, 1024)	0
add_9[0][0]		

res4c_branch2a (Conv2D)	(None, 8, 8, 256)	262400
activation_28[0][0]		
bn4c_branch2a (BatchNormalizati	(None, 8, 8, 256)	1024
res4c_branch2a[0][0]		
activation_29 (Activation)	(None, 8, 8, 256)	0
bn4c_branch2a[0][0]		
res4c_branch2b (Conv2D)	(None, 8, 8, 256)	590080
activation_29[0][0]		
bn4c_branch2b (BatchNormalizati	(None, 8, 8, 256)	1024
res4c_branch2b[0][0]		
activation_30 (Activation)	(None, 8, 8, 256)	0
bn4c_branch2b[0][0]		
res4c_branch2c (Conv2D)	(None, 8, 8, 1024)	263168
activation_30[0][0]		
bn4c_branch2c (BatchNormalizati	(None, 8, 8, 1024)	4096
res4c_branch2c[0][0]		
add_10 (Add)	(None, 8, 8, 1024)	0
bn4c_branch2c[0][0]		
activation_28[0][0]		
activation_31 (Activation)	(None, 8, 8, 1024)	0
add_10[0][0]		
res4d_branch2a (Conv2D)	(None, 8, 8, 256)	262400
activation_31[0][0]		
bn4d_branch2a (BatchNormalizati	(None, 8, 8, 256)	1024
res4d_branch2a[0][0]		
activation_32 (Activation)	(None, 8, 8, 256)	0
bn4d_branch2a[0][0]		
res4d_branch2b (Conv2D)	(None, 8, 8, 256)	590080
activation_32[0][0]		
bn4d_branch2b (BatchNormalizati	(None, 8, 8, 256)	1024
res4d_branch2b[0][0]		
activation_33 (Activation)	(None, 8, 8, 256)	0
bn4d_branch2b[0][0]		

res4d_branch2c (Conv2D) (None, 8, 8, 1024) 263168
activation_33[0][0]

bn4d_branch2c (BatchNormalizati (None, 8, 8, 1024) 4096
res4d_branch2c[0][0]

add_11 (Add) (None, 8, 8, 1024) 0
bn4d_branch2c[0][0]

activation_31[0][0]

activation_34 (Activation) (None, 8, 8, 1024) 0
add_11[0][0]

res4e_branch2a (Conv2D) (None, 8, 8, 256) 262400
activation_34[0][0]

bn4e_branch2a (BatchNormalizati (None, 8, 8, 256) 1024
res4e_branch2a[0][0]

activation_35 (Activation) (None, 8, 8, 256) 0
bn4e_branch2a[0][0]

res4e_branch2b (Conv2D) (None, 8, 8, 256) 590080
activation_35[0][0]

bn4e_branch2b (BatchNormalizati (None, 8, 8, 256) 1024
res4e_branch2b[0][0]

activation_36 (Activation) (None, 8, 8, 256) 0
bn4e_branch2b[0][0]

res4e_branch2c (Conv2D) (None, 8, 8, 1024) 263168
activation_36[0][0]

bn4e_branch2c (BatchNormalizati (None, 8, 8, 1024) 4096
res4e_branch2c[0][0]

add_12 (Add) (None, 8, 8, 1024) 0
bn4e_branch2c[0][0]

activation_34[0][0]

activation_37 (Activation) (None, 8, 8, 1024) 0
add_12[0][0]

res4f_branch2a (Conv2D) (None, 8, 8, 256) 262400
activation_37[0][0]

bn4f_branch2a (BatchNormalizati (None, 8, 8, 256) 1024

res4f_branch2a[0][0]

activation_38 (Activation) (None, 8, 8, 256) 0
bn4f_branch2a[0][0]

res4f_branch2b (Conv2D) (None, 8, 8, 256) 590080
activation_38[0][0]

bn4f_branch2b (BatchNormalizati (None, 8, 8, 256) 1024
res4f_branch2b[0][0]

activation_39 (Activation) (None, 8, 8, 256) 0
bn4f_branch2b[0][0]

res4f_branch2c (Conv2D) (None, 8, 8, 1024) 263168
activation_39[0][0]

bn4f_branch2c (BatchNormalizati (None, 8, 8, 1024) 4096
res4f_branch2c[0][0]

add_13 (Add) (None, 8, 8, 1024) 0
bn4f_branch2c[0][0]

activation_37[0][0]

activation_40 (Activation) (None, 8, 8, 1024) 0
add_13[0][0]

res5a_branch2a (Conv2D) (None, 4, 4, 512) 524800
activation_40[0][0]

bn5a_branch2a (BatchNormalizati (None, 4, 4, 512) 2048
res5a_branch2a[0][0]

activation_41 (Activation) (None, 4, 4, 512) 0
bn5a_branch2a[0][0]

res5a_branch2b (Conv2D) (None, 4, 4, 512) 2359808
activation_41[0][0]

bn5a_branch2b (BatchNormalizati (None, 4, 4, 512) 2048
res5a_branch2b[0][0]

activation_42 (Activation) (None, 4, 4, 512) 0
bn5a_branch2b[0][0]

res5a_branch2c (Conv2D) (None, 4, 4, 2048) 1050624
activation_42[0][0]

res5a_branch1 (Conv2D) (None, 4, 4, 2048) 2099200
activation_40[0][0]

bn5a_branch2c (BatchNormalizati (None, 4, 4, 2048) 8192
res5a_branch2c[0][0]

bn5a_branch1 (BatchNormalizatio (None, 4, 4, 2048) 8192
res5a_branch1[0][0]

add_14 (Add) (None, 4, 4, 2048) 0
bn5a_branch2c[0][0]

bn5a_branch1[0][0]

activation_43 (Activation) (None, 4, 4, 2048) 0
add_14[0][0]

res5b_branch2a (Conv2D) (None, 4, 4, 512) 1049088
activation_43[0][0]

bn5b_branch2a (BatchNormalizati (None, 4, 4, 512) 2048
res5b_branch2a[0][0]

activation_44 (Activation) (None, 4, 4, 512) 0
bn5b_branch2a[0][0]

res5b_branch2b (Conv2D) (None, 4, 4, 512) 2359808
activation_44[0][0]

bn5b_branch2b (BatchNormalizati (None, 4, 4, 512) 2048
res5b_branch2b[0][0]

activation_45 (Activation) (None, 4, 4, 512) 0
bn5b_branch2b[0][0]

res5b_branch2c (Conv2D) (None, 4, 4, 2048) 1050624
activation_45[0][0]

bn5b_branch2c (BatchNormalizati (None, 4, 4, 2048) 8192
res5b_branch2c[0][0]

add_15 (Add) (None, 4, 4, 2048) 0
bn5b_branch2c[0][0]

activation_43[0][0]

activation_46 (Activation) (None, 4, 4, 2048) 0
add_15[0][0]

res5c_branch2a (Conv2D) (None, 4, 4, 512) 1049088
activation_46[0][0]

```
bn5c_branch2a (BatchNormalizati (None, 4, 4, 512)      2048
    res5c_branch2a[0][0]

-----
activation_47 (Activation)      (None, 4, 4, 512)      0
    bn5c_branch2a[0][0]

-----
res5c_branch2b (Conv2D)        (None, 4, 4, 512)      2359808
    activation_47[0][0]

-----
bn5c_branch2b (BatchNormalizati (None, 4, 4, 512)      2048
    res5c_branch2b[0][0]

-----
activation_48 (Activation)      (None, 4, 4, 512)      0
    bn5c_branch2b[0][0]

-----
res5c_branch2c (Conv2D)        (None, 4, 4, 2048)     1050624
    activation_48[0][0]

-----
bn5c_branch2c (BatchNormalizati (None, 4, 4, 2048)     8192
    res5c_branch2c[0][0]

-----
add_16 (Add)                  (None, 4, 4, 2048)     0
    bn5c_branch2c[0][0]

    activation_46[0][0]

-----
activation_49 (Activation)      (None, 4, 4, 2048)     0
    add_16[0][0]

-----
avg_pool (GlobalAveragePooling2 (None, 2048)           0
    activation_49[0][0]

-----
dropout_1 (Dropout)            (None, 2048)           0
    avg_pool[0][0]

-----
dense_1 (Dense)                (None, 5004)          10253196
    dropout_1[0][0]
=====
=====

Total params: 33,840,908
Trainable params: 33,787,788
Non-trainable params: 53,120
```

```
In [12]:  
gc.collect()
```

```
Out[12]:  
0
```

Load model and weights from disc

In [13]:

```
# from keras.models import load_model

# # returns a compiled model
# # identical to the previous cell
# model = load_model(MODEL)
# print("Loaded model architecture from disk")
# gc.collect()

# model.load_weights(WEIGHTS)
# print("Loaded model weights from disk")
# model.summary()

# gc.collect()
```

Train with ImageDataGenerator and flow_from_dataframe

In [14]:

```
from keras.callbacks import LambdaCallback, ModelCheckpoint

ROTATE = 20
SEED = 28
BATCH_SIZE = 100

gc.collect()

batch_gc_callback = LambdaCallback(
    on_epoch_begin=lambda epoch, logs: gc.collect())

checkpointer = ModelCheckpoint(filepath='weights.hdf5',
                               verbose=1, save_best_only=True
)

train_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
    rescale=1./255,
    fill_mode='nearest'
)

train_generator = train_datagen.flow_from_dataframe(
    dataframe=whales_train,
    subset = "training",
    directory=TRAIN_CROPPED_IN,
    x_col="Image",
    y_col="Id",
    has_ext=True,
    seed = SEED,
    color_mode= "rgb",
    target_size=(IMAGE_HEIGHT, IMAGE_WIDTH),
    batch_size=BATCH_SIZE,
    class_mode='categorical')

gc.collect()
```

Found 15697 images belonging to 5004 classes.

Out[14]:

Visualize augmented data

In [15]:

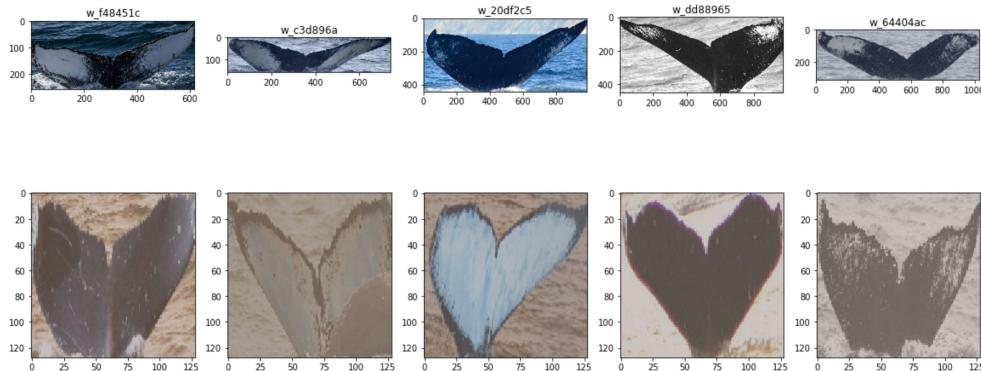
```
from skimage.io import imread
import PIL.Image as pimage

def plot_images(images_names, path):
    fig, m_axs = plt.subplots(1, len(images), figsize = (20, 10))
    #show the images and label them
    for ii, c_ax in enumerate(m_axs):
        img = imread(os.path.join(path,images_names[ii][0]))
        c_ax.imshow(img)
        c_ax.set_title(images_names[ii][1])

def plot_loaded_images(images_loaded, labels):
    fig, m_axs = plt.subplots(1, len(images_loaded), figsize = (20, 10))
    #show the images and label them
    for ii, c_ax in enumerate(m_axs):
        img = pimage.fromarray(images_loaded[ii], "RGB")
        c_ax.imshow((images_loaded[ii] + 1) / 2)
    #        c_ax.set_title(labels[ii])

#get the first 5 whale images
images = [(whale_img, whale_label) for (whale_img, whale_label) in zip(whales_train.Image[:5], whales_train.Id[:5])]
plot_images(images, TRAIN_CROPPED_IN)

x_batch, y_batch = next(train_generator)
plot_loaded_images(x_batch[:5], y_batch[:5])
```



Learning rate finder

In [16]:

```
import keras.backend as K
from keras.callbacks import Callback

class LRFinder(Callback):
    ...
    A simple callback for finding the optimal learning rate range for your model + dataset.

    # Usage
    ````python
 lr_finder = LRFinder(min_lr=1e-5,
 max_lr=1e-2,
 steps_per_epoch=np.ceil(epoch
 _size/batch_size),
 epochs=3)
```

```

 model.fit(X_train, Y_train, callbacks=[lr_finder])

 lr_finder.plot_loss()
 ...

Arguments
 min_lr: The lower bound of the learning rate range for
 the experiment.
 max_lr: The upper bound of the learning rate range for
 the experiment.
 steps_per_epoch: Number of mini-batches in the database. Calculated as `np.ceil(epoch_size/batch_size)`.

 epochs: Number of epochs to run experiment. Usually between 2 and 4 epochs is sufficient.

References
 Blog post: jeremyjordan.me/nn-learning-rate
 Original paper: https://arxiv.org/abs/1506.01186
 ...

def __init__(self, min_lr=1e-5, max_lr=1e-1, steps_per_epoch=None, epochs=None, beta=.98):
 super().__init__()

 self.min_lr = min_lr
 self.max_lr = max_lr
 self.total_iterations = steps_per_epoch * epochs
 self.iteration = 0
 self.history = {}
 self.beta = beta
 self.lr_mult = (max_lr/min_lr)**(1/steps_per_epoch)

def clr(self):
 '''Calculate the learning rate.'''
 x = self.iteration / self.total_iterations
 return self.min_lr * (self.lr_mult**self.iteration)

def on_train_begin(self, logs=None):
 '''Initialize the learning rate to the minimum value at the start of training.'''
 self.best_loss = 1e9
 self.avg_loss = 0
 self.losses, self.smoothed_losses, self.lrs, self.iterations = [], [], [], []
 logs = logs or {}
 K.set_value(self.model.optimizer.lr, self.min_lr)

def on_epoch_end(self, epoch, logs=None):
 print("Epoch ended learning rate: " + str(K.get_value(self.model.optimizer.lr)))

def on_epoch_begin(self, epoch, logs=None):
 print("Epoch start learning rate: " + str(K.get_value(self.model.optimizer.lr)))

def on_batch_end(self, epoch, logs=None):
 '''Record previous batch statistics and update the learning rate.'''
 logs = logs or {}
 self.iteration += 1

 #addition
 loss = logs.get('loss')
 self.avg_loss = self.beta * self.avg_loss + (1 - self.beta) * loss
 smoothed_loss = self.avg_loss / (1 - self.beta**self.

```

```

iteration)
 # Check if the loss is not exploding
 if self.iteration>1 and smoothed_loss > self.best_loss * 1.5:
 self.model.stop_training = True
 return
 if smoothed_loss < self.best_loss or self.iteration==1:
 self.best_loss = smoothed_loss

 lr = self.clr()
 self.losses.append(loss)
 self.smoothed_losses.append(smoothed_loss)
 self.lrs.append(lr)
 self.iterations.append(self.iteration)

 self.history.setdefault('lr', []).append(K.get_value(
 self.model.optimizer.lr))
 self.history.setdefault('iterations', []).append(self
 .iteration)

 for k, v in logs.items():
 self.history.setdefault(k, []).append(v)

 K.set_value(self.model.optimizer.lr, lr)

def plot_lr(self):
'''Helper function to quickly inspect the learning rate schedule.'''
plt.plot(self.history['iterations'], self.history['l
r'])
plt.yscale('log')
plt.xlabel('Iteration')
plt.ylabel('Learning rate')
plt.show()

def plot_lr(self):
 plt.xlabel('Iterations')
 plt.ylabel('Learning rate')
 plt.plot(self.iterations, self.lrs)
 plt.yscale('log')
 plt.show()

def plot_loss(self):
 '''Helper function to quickly observe the learning rate experiment results.'''
 plt.plot(self.history['lr'], self.history['loss'])
 plt.xscale('log')
 plt.xlabel('Learning rate')
 plt.ylabel('Loss')
 plt.show()

def plot(self, n_skip=10):
 plt.ylabel('Loss')
 plt.xlabel('Learning rate (log scale)')
 plt.plot(self.lrs[n_skip:-5], self.losses[n_skip:-5])
 plt.xscale('log')

def plot_loss2(self):
 plt.ylabel('Losses')
 plt.xlabel('Iterations')
 plt.plot(self.iterations[10:], self.losses[10:])

def plot_smoothed_loss(self, n_skip=6):
 plt.ylabel('Smoothed Losses')
 plt.xlabel('Learning rate (log scale)')
 plt.plot(self.lrs[n_skip:5], self.smoothed_losses[n

```

```

 plt.plot(self.lrs[n_skip:-5], self.smoothed_losses[n_
skip:-5])
 plt.xscale('log')

 def plot_loss_change(self, sma=1, n_skip=20, y_lim=(-0.01
, 0.01), should_lim_y = False):
 """
 Plots rate of change of the loss function.

 Parameters:
 sched - learning rate scheduler, an instance of LR
 _Finder class.
 sma - number of batches for simple moving average
 to smooth out the curve.
 n_skip - number of batches to skip on the left.
 y_lim - limits for the y axis.
 """
 derivatives = [0] * (sma + 1)
 for i in range(1 + sma, len(self.lrs)):
 derivative = (self.losses[i] - self.losses[i - sm
a]) / sma
 derivatives.append(derivative)

 plt.ylabel("d/loss")
 plt.xlabel("learning rate (log scale)")
 plt.plot(self.lrs[n_skip:], derivatives[n_skip:])
 plt.xscale('log')
 if should_lim_y:
 plt.ylim(y_lim)

```

## Train

In [17]:

### lr finder

Python notebook using data from [multiple data sources](#) • 32 views • multiple data sources Edit tags

0

Edit

```

lr_finder = LRFinder(min_lr=1e-6,
 max_lr=1,
 steps_per_epoch=np.ceil(train_generator.
n//train_generator.batch_size),
 epochs=EPOCHS)

Train the loaded model for more epochs
history = model.fit_generator(generator=train_generator,
 steps_per_epoch=STEP_SIZE_TRAIN
 ,
 epochs=EPOCHS,
 callbacks=[batch_gc_callback, lr
_finder])

```

### Version 9

10 commits

forked from Xception

### Notebook

Data

Output

Log

Comments

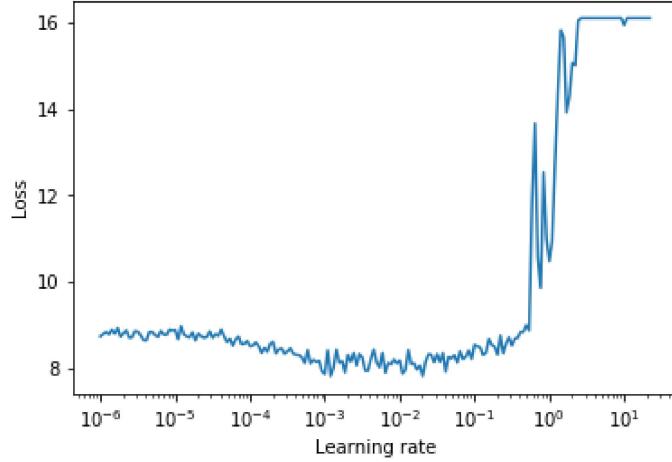
```

Epoch 1/5
Epoch start learning rate: 1e-06
156/156 [=====] - 165s 1s/step - loss: 8.5787 - acc: 0.0049 - mean_absolute_error: 3.9936e-04 - top_5_accuracy: 0.0171
Epoch ended learning rate: 1.0
Epoch 2/5
Epoch start learning rate: 1.0
37/156 [=====>.....] - ETA: 1:47 - loss: 15.4974 - acc: 2.7027e-04 - mean_absolute_error: 3.9955e-04 - top_5_accuracy: 0.7973 Epoch ended learning rate: 24.24462

```

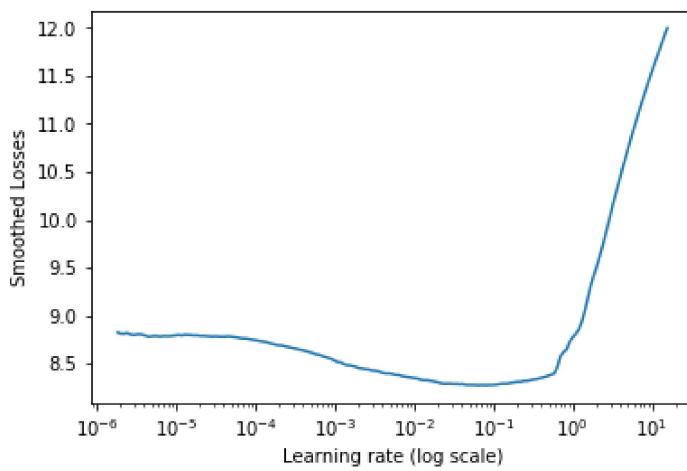
In [18]:

```
lr_finder.plot_loss()
```



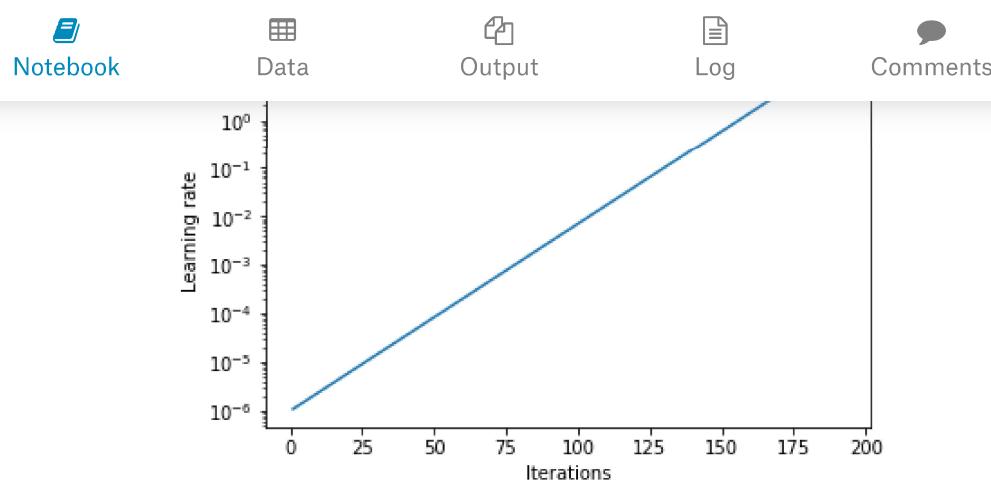
In [19]:

```
lr_finder.plot_smoothed_loss()
```



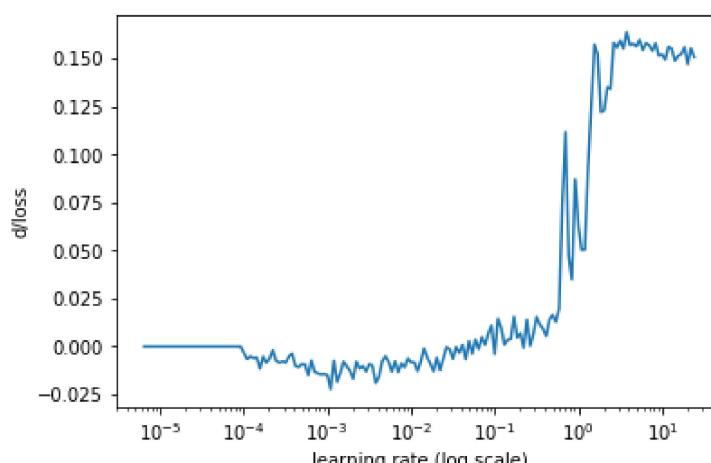
In [20]:

```
lr_finder.plot_lr()
```

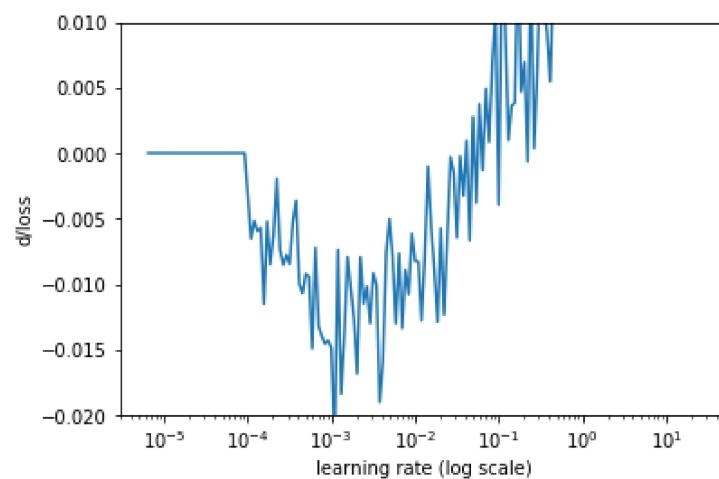


In [21]:

```
lr_finder.plot_loss_change(sma=50)
```



```
In [22]:
 lr_finder.plot_loss_change(sma=50, y_lim=(-0.02, 0.01), should
 _lim_y=True)
```



## Plot train results

```
In [23]:
def plot_accuracy(history, should_plot_val = False, should_pl
ot_top5 = False):
 acc = history.history['acc']
 l1 = plt.plot(acc, label='acc')

 if should_plot_val:
 val_acc = history.history['val_acc']
 l2 = plt.plot(val_acc, label='val_acc')

 if should_plot_top5:
 top5_acc = history.history['top_5_accuracy']
 l3 = plt.plot(top5_acc, label='top_5_accuracy')

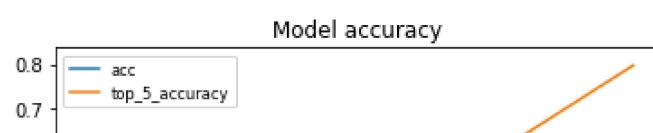
 plt.legend(loc=2, fontsize="small")
 plt.title('Model accuracy')
 plt.ylabel('Accuracy')
 plt.xlabel('Epoch')
 plt.show()

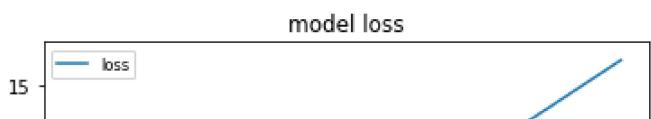
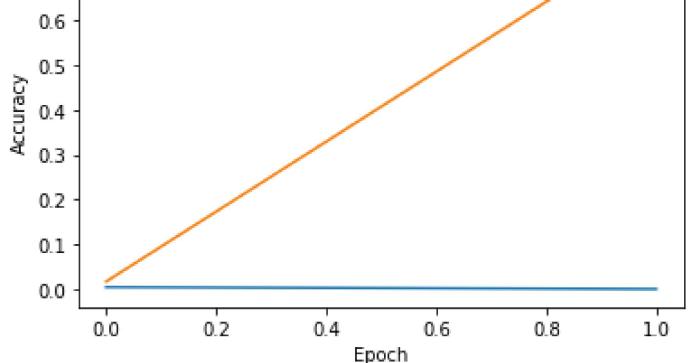
def plot_loss(history, should_plot_val = False):
 loss = history.history['loss']
 l1 = plt.plot(loss, label='loss')

 if should_plot_val:
 val_loss = history.history['val_loss']
 plt.plot(val_loss, label='val_loss')

 plt.legend(loc=2, fontsize="small")
 plt.title('model loss')
 plt.ylabel('loss')
 plt.xlabel('epoch')
 plt.show()
```

```
In [24]:
plot_accuracy(history, False, True)
plot_loss(history, False)
```





This kernel has been released under the [Apache 2.0](#) open source license.

Did you find this Kernel useful?  
Show your appreciation with an upvote

0

## Data

### Data Sources

- ▼ 🏆 Humpback Whale I...
  - grid sa... 7960 x 2
  - grid trai... 25.4k x 2
- ▼ 📄 test.zip
  - grid 0027089a4.jpg
  - grid 00313e2d2.jpg
  - grid 004344e9f.jpg
  - grid 008a4bc86.jpg
  - grid 00ac0fcfa6.jpg
  - grid 00ff45291.jpg
  - grid 012dbdb59.jpg
  - grid 0169cec0e.jpg
  - grid 01830c9cf.jpg
  - grid 01b1ecf7b.jpg
  - ... 1000+ more
- ▼ 📄 train.zip
  - grid 002b4615d.jpg
  - grid 00600ce17.jpg
  - grid 00d641885.jpg
  - grid 00eaedfab.jpg
  - grid 00fee3975.jpg
  - grid 010a1f0eb.jpg
  - grid 01237f1ce.jpg
  - grid 01dcba20f.jpg
  - grid 0202dfb29.jpg
  - grid 020ab0f9b.jpg
  - ... 1000+ more



### Humpback Whale Identification

Can you identify a whale by its tail?

Last Updated: 2 months ago

#### About this Competition

This training data contains thousands of images of humpback whale flukes. Individual whales have been identified by researchers and given an `Id`. The challenge is to predict the whale `Id` of images in the test set. What makes this such a challenge is that there are only a few examples for each of 3,000+ whale `Ids`.

#### File descriptions

- **train.zip** - a folder containing the training images
- **train.csv** - maps the training `Image` to the appropriate whale `Id`. Whales that are not predicted to have a label identified in the training data should be labeled as `new_whale`.
- **test.zip** - a folder containing the test images to predict the whale `Id`
- **sample\_submission.csv** - a sample submission file in the correct format

## Output Files

- Model\_ResNet50.h5
- Weights\_ResNet50....

## About this file

This file was created from a Kernel, it does not have a description.

### Model\_ResNet50.h5



We don't support previews for this file yet

## Run Info

Succeeded	True	Run Time	277.9 seconds
Exit Code	0	Queue Time	0 seconds
Docker Image Name	/python(Dockerfile)	Output Size	0
Timeout Exceeded	False	Used All Space	False
Failure Message			

## Log

[Download Log](#)

```
Time Line # Log Message
2.7s 1 [NbConvertApp] Converting notebook __notebook__.ipynb to
 notebook
2.8s 2 [NbConvertApp] Executing notebook with kernel: python3
7.4s 3 2019-02-02 20:46:15.098146: I
 tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:964]
 successful NUMA node read from SysFS had negative value (-1),
 but there must be at least one NUMA node, so returning NUMA
 node zero
7.4s 4 2019-02-02 20:46:15.104368: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:1432] Found
 device 0 with properties:
 name: Tesla K80 major: 3 minor: 7 memoryClockRate(GHz): 0.8235
 pciBusID: 0000:00:04.0
 totalMemory: 11.17GiB freeMemory: 11.10GiB
 2019-02-02 20:46:15.104406: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:1511] Adding
 visible gpu devices: 0
7.9s 5 2019-02-02 20:46:15.573820: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:982] Device
 interconnect StreamExecutor with strength 1 edge matrix:
7.9s 6 2019-02-02 20:46:15.573966: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:988] 0
 2019-02-02 20:46:15.574070: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:1001] 0: N
 2019-02-02 20:46:15.574559: I
 tensorflow/core/common_runtime/gpu/gpu_device.cc:1115] Created
 TensorFlow device
 (/job:localhost/replica:0/task:0/device:GPU:0 with 10758 MB
 memory) -> physical GPU (device: 0, name: Tesla K80, pci bus
 id: 0000:00:04.0, compute capability: 3.7)
274.7s 7 [NbConvertApp] Writing 718399 bytes to __notebook__.ipynb
276.9s 8 [NbConvertApp] Converting notebook __notebook__.ipynb to html
277.5s 9 [NbConvertApp] Support files will be in __results__files/
277.5s 10 [NbConvertApp] Making directory __results__files
 [NbConvertApp] Writing 365882 bytes to __results__.html
277.5s 11
277.5s 13 Complete. Exited with code 0.
```

## Comments (0)



Click here to enter a comment...

