



In [1]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load in

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory

import os
print(os.listdir("../input"))
import keras
print(keras.__version__)
# Any results you write to the current directory are saved as output.
```

[ 'humpback-whale-identification', 'whales-cropped' ]

Using TensorFlow backend.

## 2.2.4

In [2]:

```
HW = 'humpback-whale-identification'
# TRAIN = '../input/humpback-whale-identification/train/'
TRAIN_CROPPED = "whales-cropped/cropped_train/cropped_train/"
TRAIN_CROPPED_IN = '../input/' + TRAIN_CROPPED

# TEST = '../input/humpback-whale-identification/test/'
TEST_CROPPED = "whales-cropped/cropped_test/cropped_test/"
TEST_CROPPED_IN = '../input/' + TEST_CROPPED

LABELS = '../input/humpback-whale-identification/train.csv'
SAMPLE_SUB = '../input/humpback-whale-identification/sample_submission.csv'

train = pd.read_csv(LABELS)
print("With new_whale:")
train.head()
```

With new\_whale:

Out[2]:

	Image	Id
0	0000e88ab.jpg	w_f48451c
1	0001f9222.jpg	w_c3d896a
2	00029d126.jpg	w_20df2c5
3	00050a15a.jpg	new_whale
4	0005c1ef8.jpg	new_whale

In [3]:

```
MODEL_F = 'Model_Xception.h5'
WEIGHTS_F = 'Weights_Xception.h5'
MODEL = '../input/Xception/' + MODEL_F
WEIGHTS = '../input/Xception/' + WEIGHTS_F
```

In [4]:

```
train.describe()
```

Out[4]:

	Image	Id
count	25361	25361
unique	25361	5005
top	e5982ded0.jpg	new_whale
freq	1	9664

In [5]:

```
import random
from IPython.display import Image
print("Example whale image")

#show sample image
name = random.choice(train['Image'])
print(name)
Image(filename = TRAIN_CROPPED_IN + name)
```

Example whale image  
b576308d0.jpg

Out[5]:



In [6]:

```
criteria = train['Id'] != 'new_whale'
whales_train = train[criteria]
```

```
print("Without new_whale:")
whales_train.head()
```

Without new\_whale:

Out[6]:

	Image	Id
0	0000e88ab.jpg	w_f48451c
1	0001f9222.jpg	w_c3d896a
2	00029d126.jpg	w_20df2c5
6	000a6daec.jpg	w_dd88965
8	0016b897a.jpg	w_64404ac

In [7]:

```
unique_labels = np.unique(whales_train.Id.values)
labels_list = unique_labels
```

In [8]:

```
whales_train.describe()
```

Out[8]:

	Image	Id
count	15697	15697
unique	15697	5004
top	b4284d5ef.jpg	w_23a388d
freq	1	73

In [9]:

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from matplotlib.pyplot import imshow

from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder

from keras import layers
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator

# from keras.applications.imagenet_utils import preprocess_input
# from keras.applications.resnet50 import ResNet50, preprocess_input
from keras.applications.xception import Xception, preprocess_input
# from keras.applications.inception_resnet_v2 import InceptionResNetV2, preprocess_input

from keras.losses import binary_crossentropy
from keras.preprocessing.image import ImageDataGenerator
```

```
from keras.layers import Input, Dense, Activation,  
BatchNormalization, Flatten, Conv2D, GlobalAverageP  
ooling2D  
from keras.layers import AveragePooling2D, MaxPooli  
ng2D, Dropout  
from keras.models import Model  
from keras.metrics import top_k_categorical_accurac  
y  
  
import keras.backend as K  
from keras.models import Sequential  
from PIL import Image  
import gc  
import warnings  
warnings.simplefilter("ignore", category=Deprecatio  
nWarning)  
  
%matplotlib inline
```

In [10]:

```
IMAGE_HEIGHT = 128  
IMAGE_WIDTH = 128  
IMAGE_SHAPE = (IMAGE_HEIGHT, IMAGE_WIDTH, 3)
```

In [11]:

```
CLASSES = 5004  
EPOCHS = 5  
BATCH_SIZE = 32  
  
def top_5_accuracy(y_true, y_pred):  
    return top_k_categorical_accuracy(y_true, y_pre  
d, k=5)  
  
# setup model  
base_model = Xception(weights='imagenet', include_t  
op=False, input_shape = IMAGE_SHAPE)  
  
x = base_model.output  
x = GlobalAveragePooling2D(name='avg_pool')(x)  
x = Dropout(0.4)(x)  
predictions = Dense(CLASSES, activation='softmax')(  
x)  
model = Model(inputs=base_model.input, outputs=pred  
ictions)  
  
# transfer learning  
for layer in base_model.layers:  
    layer.trainable = True  
  
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy', 'mae', top_5_acc  
uracy])  
  
model.summary()
```

```
_weights_tf_dim_ordering_tf_kernels_notop.h5
83689472/83683744 [=====]
- 2s 0us/step

-----
-----  
Layer (type)           Output Shape  
Param #    Connected to  
=====
```

---

```
input_1 (InputLayer)      (None, 128, 128, 3)  
0
```

---

```
block1_conv1 (Conv2D)      (None, 63, 63, 32)  
864          input_1[0][0]
```

---

```
block1_conv1_bn (BatchNormaliza (None, 63, 63, 32)  
128          block1_conv1[0][0]
```

---

```
block1_conv1_act (Activation)  (None, 63, 63, 32)  
0            block1_conv1_bn[0][0]
```

---

```
block1_conv2 (Conv2D)      (None, 61, 61, 64)  
18432         block1_conv1_act[0][0]
```

---

```
block1_conv2_bn (BatchNormaliza (None, 61, 61, 64)  
256          block1_conv2[0][0]
```

---

```
block1_conv2_act (Activation)  (None, 61, 61, 64)  
0            block1_conv2_bn[0][0]
```

---

```
block2_sepconv1 (SeparableConv2 (None, 61, 61, 128)  
8768          block1_conv2_act[0][0]
```

---

```
block2_sepconv1_bn (BatchNormal (None, 61, 61, 128)  
512          block2_sepconv1[0][0]
```

---

```
block2_sepconv2_act (Activation (None, 61, 61, 128)  
0            block2_sepconv1_bn[0][0]
```

---

```
block2_sepconv2 (SeparableConv2 (None, 61, 61, 128)  
17536         block2_sepconv2_act[0][0]
```

---

```
block2_sepconv2_bn (BatchNormal (None, 61, 61, 128)  
512          block2_sepconv2[0][0]
```

---

```
conv2d_1 (Conv2D)        (None, 31, 31, 128)  
8192          block1_conv2_act[0][0]
```

```
-----  
block2_pool (MaxPooling2D)      (None, 31, 31, 128)  
0          block2_sepconv2_bn[0][0]  
-----  
  
batch_normalization_1 (BatchNor (None, 31, 31, 128)  
512        conv2d_1[0][0]  
-----  
  
add_1 (Add)                  (None, 31, 31, 128)  
0          block2_pool[0][0]  
  
batch_normalization_1[0][0]  
-----  
  
block3_sepconv1_act (Activation (None, 31, 31, 128)  
0          add_1[0][0]  
-----  
  
block3_sepconv1 (SeparableConv2 (None, 31, 31, 256)  
33920      block3_sepconv1_act[0][0]  
-----  
  
block3_sepconv1_bn (BatchNormal (None, 31, 31, 256)  
1024       block3_sepconv1[0][0]  
-----  
  
block3_sepconv2_act (Activation (None, 31, 31, 256)  
0          block3_sepconv1_bn[0][0]  
-----  
  
block3_sepconv2 (SeparableConv2 (None, 31, 31, 256)  
67840      block3_sepconv2_act[0][0]  
-----  
  
block3_sepconv2_bn (BatchNormal (None, 31, 31, 256)  
1024       block3_sepconv2[0][0]  
-----  
  
conv2d_2 (Conv2D)            (None, 16, 16, 256)  
32768      add_1[0][0]  
-----  
  
block3_pool (MaxPooling2D)    (None, 16, 16, 256)  
0          block3_sepconv2_bn[0][0]  
-----  
  
batch_normalization_2 (BatchNor (None, 16, 16, 256)  
1024       conv2d_2[0][0]  
-----  
  
add_2 (Add)                  (None, 16, 16, 256)  
0          block3_pool[0][0]  
  
batch_normalization_2[0][0]  
-----  
  
block4_sepconv1_act (Activation (None, 16, 16, 256)  
0          add_2[0][0]
```

-----  
block4\_sepconv1 (SeparableConv2 (None, 16, 16, 728)  
188672 block4\_sepconv1\_act[0][0]  
-----  
block4\_sepconv1\_bn (BatchNormal (None, 16, 16, 728)  
2912 block4\_sepconv1[0][0]  
-----  
block4\_sepconv2\_act (Activation (None, 16, 16, 728)  
0 block4\_sepconv1\_bn[0][0]  
-----  
block4\_sepconv2 (SeparableConv2 (None, 16, 16, 728)  
536536 block4\_sepconv2\_act[0][0]  
-----  
block4\_sepconv2\_bn (BatchNormal (None, 16, 16, 728)  
2912 block4\_sepconv2[0][0]  
-----  
conv2d\_3 (Conv2D) (None, 8, 8, 728)  
186368 add\_2[0][0]  
-----  
block4\_pool (MaxPooling2D) (None, 8, 8, 728)  
0 block4\_sepconv2\_bn[0][0]  
-----  
batch\_normalization\_3 (BatchNor (None, 8, 8, 728)  
2912 conv2d\_3[0][0]  
-----  
add\_3 (Add) (None, 8, 8, 728)  
0 block4\_pool[0][0]  
batch\_normalization\_3[0][0]  
-----  
block5\_sepconv1\_act (Activation (None, 8, 8, 728)  
0 add\_3[0][0]  
-----  
block5\_sepconv1 (SeparableConv2 (None, 8, 8, 728)  
536536 block5\_sepconv1\_act[0][0]  
-----  
block5\_sepconv1\_bn (BatchNormal (None, 8, 8, 728)  
2912 block5\_sepconv1[0][0]  
-----  
block5\_sepconv2\_act (Activation (None, 8, 8, 728)  
0 block5\_sepconv1\_bn[0][0]  
-----  
block5\_sepconv2 (SeparableConv2 (None, 8, 8, 728)  
536536 block5\_sepconv2\_act[0][0]  
-----

```
block5_sepconv2_bn (BatchNormal (None, 8, 8, 728)
2912          block5_sepconv2[0][0]
-----
-----  
block5_sepconv3_act (Activation (None, 8, 8, 728)
0          block5_sepconv2_bn[0][0]
-----  
block5_sepconv3 (SeparableConv2 (None, 8, 8, 728)
536536          block5_sepconv3_act[0][0]
-----  
block5_sepconv3_bn (BatchNormal (None, 8, 8, 728)
2912          block5_sepconv3[0][0]
-----  
add_4 (Add)          (None, 8, 8, 728)
0          block5_sepconv3_bn[0][0]
-----  
add_3[0][0]
-----  
block6_sepconv1_act (Activation (None, 8, 8, 728)
0          add_4[0][0]
-----  
block6_sepconv1 (SeparableConv2 (None, 8, 8, 728)
536536          block6_sepconv1_act[0][0]
-----  
block6_sepconv1_bn (BatchNormal (None, 8, 8, 728)
2912          block6_sepconv1[0][0]
-----  
block6_sepconv2_act (Activation (None, 8, 8, 728)
0          block6_sepconv1_bn[0][0]
-----  
block6_sepconv2 (SeparableConv2 (None, 8, 8, 728)
536536          block6_sepconv2_act[0][0]
-----  
block6_sepconv2_bn (BatchNormal (None, 8, 8, 728)
2912          block6_sepconv2[0][0]
-----  
block6_sepconv3_act (Activation (None, 8, 8, 728)
0          block6_sepconv2_bn[0][0]
-----  
block6_sepconv3 (SeparableConv2 (None, 8, 8, 728)
536536          block6_sepconv3_act[0][0]
-----  
block6_sepconv3_bn (BatchNormal (None, 8, 8, 728)
2912          block6_sepconv3[0][0]
-----  
add_5 (Add)          (None, 8, 8, 728)
0          block6_sepconv3_bn[0][0]
```

add\_4[0][0]  
-----  
block7\_sepconv1\_act (Activation (None, 8, 8, 728)  
0 add\_5[0][0]  
-----  
block7\_sepconv1 (SeparableConv2 (None, 8, 8, 728)  
536536 block7\_sepconv1\_act[0][0]  
-----  
block7\_sepconv1\_bn (BatchNormal (None, 8, 8, 728)  
2912 block7\_sepconv1[0][0]  
-----  
block7\_sepconv2\_act (Activation (None, 8, 8, 728)  
0 block7\_sepconv1\_bn[0][0]  
-----  
block7\_sepconv2 (SeparableConv2 (None, 8, 8, 728)  
536536 block7\_sepconv2\_act[0][0]  
-----  
block7\_sepconv2\_bn (BatchNormal (None, 8, 8, 728)  
2912 block7\_sepconv2[0][0]  
-----  
block7\_sepconv3\_act (Activation (None, 8, 8, 728)  
0 block7\_sepconv2\_bn[0][0]  
-----  
block7\_sepconv3 (SeparableConv2 (None, 8, 8, 728)  
536536 block7\_sepconv3\_act[0][0]  
-----  
block7\_sepconv3\_bn (BatchNormal (None, 8, 8, 728)  
2912 block7\_sepconv3[0][0]  
-----  
add\_6 (Add) (None, 8, 8, 728)  
0 block7\_sepconv3\_bn[0][0]  
-----  
add\_5[0][0]  
-----  
block8\_sepconv1\_act (Activation (None, 8, 8, 728)  
0 add\_6[0][0]  
-----  
block8\_sepconv1 (SeparableConv2 (None, 8, 8, 728)  
536536 block8\_sepconv1\_act[0][0]  
-----  
block8\_sepconv1\_bn (BatchNormal (None, 8, 8, 728)  
2912 block8\_sepconv1[0][0]  
-----  
block8\_sepconv2\_act (Activation (None, 8, 8, 728)  
0 block8\_sepconv1\_bn[0][0]

```
-----  
-----  
block8_sepconv2 (SeparableConv2 (None, 8, 8, 728)  
    536536      block8_sepconv2_act[0][0]  
-----  
-----  
block8_sepconv2_bn (BatchNormal (None, 8, 8, 728)  
    2912        block8_sepconv2[0][0]  
-----  
-----  
block8_sepconv3_act (Activation (None, 8, 8, 728)  
    0            block8_sepconv2_bn[0][0]  
-----  
-----  
block8_sepconv3 (SeparableConv2 (None, 8, 8, 728)  
    536536      block8_sepconv3_act[0][0]  
-----  
-----  
block8_sepconv3_bn (BatchNormal (None, 8, 8, 728)  
    2912        block8_sepconv3[0][0]  
-----  
-----  
add_7 (Add)           (None, 8, 8, 728)  
    0            block8_sepconv3_bn[0][0]  
  
                                add_6[0][0]  
-----  
-----  
block9_sepconv1_act (Activation (None, 8, 8, 728)  
    0            add_7[0][0]  
-----  
-----  
block9_sepconv1 (SeparableConv2 (None, 8, 8, 728)  
    536536      block9_sepconv1_act[0][0]  
-----  
-----  
block9_sepconv1_bn (BatchNormal (None, 8, 8, 728)  
    2912        block9_sepconv1[0][0]  
-----  
-----  
block9_sepconv2_act (Activation (None, 8, 8, 728)  
    0            block9_sepconv1_bn[0][0]  
-----  
-----  
block9_sepconv2 (SeparableConv2 (None, 8, 8, 728)  
    536536      block9_sepconv2_act[0][0]  
-----  
-----  
block9_sepconv2_bn (BatchNormal (None, 8, 8, 728)  
    2912        block9_sepconv2[0][0]  
-----  
-----  
block9_sepconv3_act (Activation (None, 8, 8, 728)  
    0            block9_sepconv2_bn[0][0]  
-----  
-----  
block9_sepconv3 (SeparableConv2 (None, 8, 8, 728)  
    536536      block9_sepconv3_act[0][0]  
-----
```

```
block9_sepconv3_bn (BatchNormal (None, 8, 8, 728)
2912          block9_sepconv3[0][0]
-----
add_8 (Add)           (None, 8, 8, 728)
0          block9_sepconv3_bn[0][0]
add_7[0][0]
-----
block10_sepconv1_act (Activatio (None, 8, 8, 728)
0          add_8[0][0]
-----
block10_sepconv1 (SeparableConv (None, 8, 8, 728)
536536      block10_sepconv1_act[0][0]
-----
block10_sepconv1_bn (BatchNorma (None, 8, 8, 728)
2912          block10_sepconv1[0][0]
-----
block10_sepconv2_act (Activatio (None, 8, 8, 728)
0          block10_sepconv1_bn[0][0]
-----
block10_sepconv2 (SeparableConv (None, 8, 8, 728)
536536      block10_sepconv2_act[0][0]
-----
block10_sepconv2_bn (BatchNorma (None, 8, 8, 728)
2912          block10_sepconv2[0][0]
-----
block10_sepconv3_act (Activatio (None, 8, 8, 728)
0          block10_sepconv2_bn[0][0]
-----
block10_sepconv3 (SeparableConv (None, 8, 8, 728)
536536      block10_sepconv3_act[0][0]
-----
block10_sepconv3_bn (BatchNorma (None, 8, 8, 728)
2912          block10_sepconv3[0][0]
-----
add_9 (Add)           (None, 8, 8, 728)
0          block10_sepconv3_bn[0][0]
add_8[0][0]
-----
block11_sepconv1_act (Activatio (None, 8, 8, 728)
0          add_9[0][0]
-----
block11_sepconv1 (SeparableConv (None, 8, 8, 728)
536536      block11_sepconv1_act[0][0]
```

```
block11_sepconv1_bn (BatchNorma (None, 8, 8, 728)
2912          block11_sepconv1[0][0]
-----
block11_sepconv2_act (Activatio (None, 8, 8, 728)
0          block11_sepconv1_bn[0][0]
-----
block11_sepconv2 (SeparableConv (None, 8, 8, 728)
536536          block11_sepconv2_act[0][0]
-----
block11_sepconv2_bn (BatchNorma (None, 8, 8, 728)
2912          block11_sepconv2[0][0]
-----
block11_sepconv3_act (Activatio (None, 8, 8, 728)
0          block11_sepconv2_bn[0][0]
-----
block11_sepconv3 (SeparableConv (None, 8, 8, 728)
536536          block11_sepconv3_act[0][0]
-----
block11_sepconv3_bn (BatchNorma (None, 8, 8, 728)
2912          block11_sepconv3[0][0]
-----
add_10 (Add)          (None, 8, 8, 728)
0          block11_sepconv3_bn[0][0]
add_9[0][0]
-----
block12_sepconv1_act (Activatio (None, 8, 8, 728)
0          add_10[0][0]
-----
block12_sepconv1 (SeparableConv (None, 8, 8, 728)
536536          block12_sepconv1_act[0][0]
-----
block12_sepconv1_bn (BatchNorma (None, 8, 8, 728)
2912          block12_sepconv1[0][0]
-----
block12_sepconv2_act (Activatio (None, 8, 8, 728)
0          block12_sepconv1_bn[0][0]
-----
block12_sepconv2 (SeparableConv (None, 8, 8, 728)
536536          block12_sepconv2_act[0][0]
-----
block12_sepconv2_bn (BatchNorma (None, 8, 8, 728)
2912          block12_sepconv2[0][0]
-----
block12_sepconv3_act (Activatio (None, 8, 8, 728)
0          block12_sepconv2_bn[0][0]
```

0 block12\_sepconv2\_bn[0][0]

-----

block12\_sepconv3 (SeparableConv (None, 8, 8, 728)  
536536 block12\_sepconv3\_act[0][0]

-----

block12\_sepconv3\_bn (BatchNorma (None, 8, 8, 728)  
2912 block12\_sepconv3[0][0]

-----

add\_11 (Add) (None, 8, 8, 728)  
0 block12\_sepconv3\_bn[0][0]

-----

add\_10[0][0]

-----

block13\_sepconv1\_act (Activatio (None, 8, 8, 728)  
0 add\_11[0][0]

-----

block13\_sepconv1 (SeparableConv (None, 8, 8, 728)  
536536 block13\_sepconv1\_act[0][0]

-----

block13\_sepconv1\_bn (BatchNorma (None, 8, 8, 728)  
2912 block13\_sepconv1[0][0]

-----

block13\_sepconv2\_act (Activatio (None, 8, 8, 728)  
0 block13\_sepconv1\_bn[0][0]

-----

block13\_sepconv2 (SeparableConv (None, 8, 8, 1024)  
752024 block13\_sepconv2\_act[0][0]

-----

block13\_sepconv2\_bn (BatchNorma (None, 8, 8, 1024)  
4096 block13\_sepconv2[0][0]

-----

conv2d\_4 (Conv2D) (None, 4, 4, 1024)  
745472 add\_11[0][0]

-----

block13\_pool (MaxPooling2D) (None, 4, 4, 1024)  
0 block13\_sepconv2\_bn[0][0]

-----

batch\_normalization\_4 (BatchNor (None, 4, 4, 1024)  
4096 conv2d\_4[0][0]

-----

add\_12 (Add) (None, 4, 4, 1024)  
0 block13\_pool[0][0]

-----

batch\_normalization\_4[0][0]

-----

block14\_sepconv1 (SeparableConv (None, 4, 4, 1536)

```
1582080      add_12[0][0]
-----
-----  
block14_sepconv1_bn (BatchNorma (None, 4, 4, 1536)
6144          block14_sepconv1[0][0]
-----
-----  
block14_sepconv1_act (Activatio (None, 4, 4, 1536)
0            block14_sepconv1_bn[0][0]
-----
-----  
block14_sepconv2 (SeparableConv (None, 4, 4, 2048)
3159552      block14_sepconv1_act[0][0]
-----
-----  
block14_sepconv2_bn (BatchNorma (None, 4, 4, 2048)
8192        block14_sepconv2[0][0]
-----
-----  
block14_sepconv2_act (Activatio (None, 4, 4, 2048)
0            block14_sepconv2_bn[0][0]
-----
-----  
avg_pool (GlobalAveragePooling2 (None, 2048)
0            block14_sepconv2_act[0][0]
-----
-----  
dropout_1 (Dropout)           (None, 2048)
0            avg_pool[0][0]
-----
-----  
dense_1 (Dense)              (None, 5004)
10253196    dropout_1[0][0]
=====  
=====  
Total params: 31,114,676  
Trainable params: 31,060,148  
Non-trainable params: 54,528
```

```
In [12]:  
gc.collect()
```

```
Out[12]:  
0
```

## Load model and weights from disc

```
In [13]:  
# from keras.models import load_model  
  
# # returns a compiled model  
# # identical to the previous cell  
# model = load_model(MODEL)  
# print("Loaded model architecture from disk")  
# gc.collect()
```

```
# model.load_weights(WEIGHTS)
# print("Loaded model weights from disk")
# model.summary()

# gc.collect()
```

## Train with ImageDataGenerator and flow\_from\_dataframe

In [14]:

```
from keras.callbacks import LambdaCallback, ModelCheckpoint

ROTATE = 20
SEED = 28
BATCH_SIZE = 100

gc.collect()

batch_gc_callback = LambdaCallback(
    on_epoch_begin=lambda epoch, logs: gc.collect())

checkpointer = ModelCheckpoint(filepath='weights.hdf5',
                               verbose=1, save_best_only=True)

train_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
    rescale=1./255,
    fill_mode='nearest'
)

train_generator = train_datagen.flow_from_dataframe(
(
    dataframe=whales_train,
    subset = "training",
    directory=TRAIN_CROPPED_IN,
    x_col="Image",
    y_col="Id",
    has_ext=True,
    seed = SEED,
    color_mode= "rgb",
    target_size=(IMAGE_HEIGHT, IMAGE_WIDTH),
    batch_size=BATCH_SIZE,
    class_mode='categorical')

gc.collect()
```

Found 15697 images belonging to 5004 classes.

Out[14]:

## Visualize augmented data

In [15]:

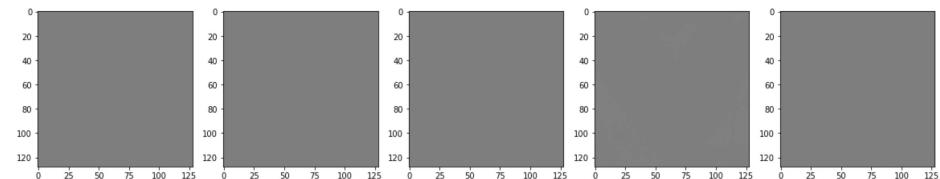
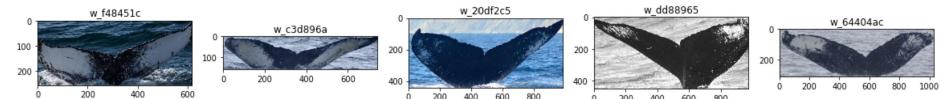
```
from skimage.io import imread
import PIL.Image as pimage

def plot_images(images_names, path):
    fig, m_axs = plt.subplots(1, len(images), figsize = (20, 10))
    #show the images and label them
    for ii, c_ax in enumerate(m_axs):
        img = imread(os.path.join(path, images_names[ii][0]))
        c_ax.imshow(img)
        c_ax.set_title(images_names[ii][1])

def plot_loaded_images(images_loaded, labels):
    fig, m_axs = plt.subplots(1, len(images_loaded), figsize = (20, 10))
    #show the images and label them
    for ii, c_ax in enumerate(m_axs):
        img = pimage.fromarray(images_loaded[ii], "RGB")
        c_ax.imshow((images_loaded[ii] + 1) / 2)
    #        c_ax.set_title(labels[ii])

#get the first 5 whale images
images = [(whale_img, whale_label) for (whale_img, whale_label) in zip(whales_train.Image[:5], whales_train.Id[:5])]
plot_images(images, TRAIN_CROPPED_IN)

x_batch, y_batch = next(train_generator)
plot_loaded_images(x_batch[:5], y_batch[:5])
```



## Learning rate finder

In [16]:

```
import keras.backend as K
from keras.callbacks import Callback
```

```
class LRFinder(Callback):
    ...
```

*A simple callback for finding the optimal learning rate range for your model + dataset.*

```
# Usage
```python
lr_finder = LRFinder(min_lr=1e-5,
                      max_lr=1e-2,
                      steps_per_epoch=np.c
eil(epoch_size/batch_size),
                      epochs=3)
model.fit(X_train, Y_train, callbacks=[lr
_finder])```
lr_finder.plot_loss()
```
```

#### **# Arguments**

*min\_lr: The lower bound of the learning rate range for the experiment.*

*max\_lr: The upper bound of the learning rate range for the experiment.*

*steps\_per\_epoch: Number of mini-batches in the dataset. Calculated as `np.ceil(epoch\_size/batch\_size)`.*

*epochs: Number of epochs to run experiment. Usually between 2 and 4 epochs is sufficient.*

#### **# References**

*Blog post: [jeremyjordan.me/nn-learning-rate](http://jeremyjordan.me/nn-learning-rate)*

*Original paper: <https://arxiv.org/abs/1506.01>*

**186**

```

```
def __init__(self, min_lr=1e-5, max_lr=1e-1, st
eps_per_epoch=None, epochs=None, beta=.98):
    super().__init__()

    self.min_lr = min_lr
    self.max_lr = max_lr
    self.total_iterations = steps_per_epoch * e
pochs
    self.iteration = 0
    self.history = {}
    self.beta = beta
    self.lr_mult = (max_lr/min_lr)**(1/steps_pe
r_epoch)

def clr(self):
    '''Calculate the learning rate.'''
    x = self.iteration / self.total_iterations
    return self.min_lr * (self.lr_mult**self.it
eration)

def on_train_begin(self, logs=None):
    '''Initialize the learning rate to the minim
um value at the start of training.'''
    self.best_loss = 1e9
    self.avg_loss = 0
    self.losses, self.smoothed_losses, self.lrs
, self.iterations = [], [], [], []
```

```

        logs = logs or {}
        K.set_value(self.model.optimizer.lr, self.m
in_lr)

        def on_epoch_end(self, epoch, logs=None):
            print(" Epoch ended learning rate: " + str(
K.get_value(self.model.optimizer.lr)))

        def on_epoch_begin(self, epoch, logs=None):
            print(" Epoch start learning rate: " + str(
K.get_value(self.model.optimizer.lr)))

        def on_batch_end(self, epoch, logs=None):
            '''Record previous batch statistics and upda
te the learning rate.'''
            logs = logs or {}
            self.iteration += 1

            #addition
            loss = logs.get('loss')
            self.avg_loss = self.beta * self.avg_loss +
(1 - self.beta) * loss
            smoothed_loss = self.avg_loss / (1 - self.b
eta**self.iteration)
            # Check if the loss is not exploding
            if self.iteration>1 and smoothed_loss > sel
f.best_loss * 1.5:
                self.model.stop_training = True
                return
            if smoothed_loss < self.best_loss or self.i
teration==1:
                self.best_loss = smoothed_loss

            lr = self.clr()
            self.losses.append(loss)
            self.smoothed_losses.append(smoothed_loss)
            self.lrs.append(lr)
            self.iterations.append(self.iteration)

            self.history.setdefault('lr', []).append(K.
get_value(self.model.optimizer.lr))
            self.history.setdefault('iterations', []).a
ppend(self.iteration)

            for k, v in logs.items():
                self.history.setdefault(k, []).append(v
)

            K.set_value(self.model.optimizer.lr,lr)

#     def plot_lr(self):
#         '''Helper function to quickly inspect the l
earning rate schedule.'''
#         plt.plot(self.history['iterations'], self.h
istory['lr'])
#         plt.yscale('log')
#         plt.xlabel('Iteration')
#         plt.ylabel('Learning rate')
#         plt.show()

```

```

        def plot_lr(self):
            plt.xlabel('Iterations')
            plt.ylabel('Learning rate')
            plt.plot(self.iterations, self.lrs)
            plt.yscale('log')
            plt.show()

        def plot_loss(self):
            '''Helper function to quickly observe the learning rate experiment results.'''
            plt.plot(self.history['lr'], self.history['loss'])
            plt.xscale('log')
            plt.xlabel('Learning rate')
            plt.ylabel('Loss')
            plt.show()

        def plot(self, n_skip=10):
            plt.ylabel('Loss')
            plt.xlabel('Learning rate (log scale)')
            plt.plot(self.lrs[n_skip:-5], self.losses[n_skip:-5])
            plt.xscale('log')

        def plot_loss2(self):
            plt.ylabel('Losses')
            plt.xlabel('Iterations')
            plt.plot(self.iterations[10:], self.losses[10:])

        def plot_smoothed_loss(self, n_skip=6):
            plt.ylabel('Smoothed Losses')
            plt.xlabel('Learning rate (log scale)')
            plt.plot(self.lrs[n_skip:-5], self.smoothed_losses[n_skip:-5])
            plt.xscale('log')

        def plot_loss_change(self, sma=1, n_skip=20, y_lim=(-0.01,0.01), should_lim_y = False):
            """
            Plots rate of change of the loss function.

            Parameters:
                sched - learning rate scheduler, an instance of LR_Finder class.
                sma - number of batches for simple moving average to smooth out the curve.
                n_skip - number of batches to skip on the left.
                y_lim - limits for the y axis.
            """
            derivatives = [0] * (sma + 1)
            for i in range(1 + sma, len(self.lrs)):
                derivative = (self.losses[i] - self.losses[i - sma]) / sma
                derivatives.append(derivative)

            plt.ylabel("d/loss")
            plt.xlabel("learning rate (log scale)")
            plt.plot(self.lrs[n_skip:], derivatives[n_skip:])

```

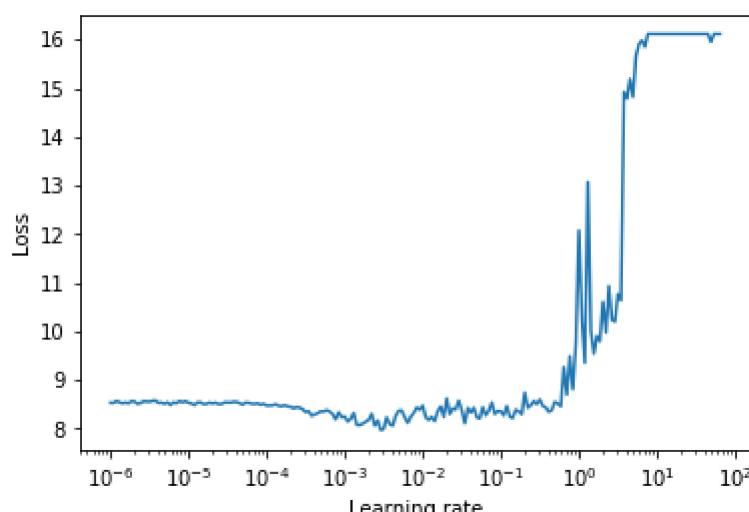
```
    plt.xscale('log')
    if should_lim_y:
        plt.ylim(y_lim)
```

## Train

```
In [17]:  
# fits the model on batches with real-time data augmentation:  
STEP_SIZE_TRAIN=train_generator.n//train_generator.  
batch_size  
  
lr_finder = LRFinder(min_lr=1e-6,  
                      max_lr=1,  
                      steps_per_epoch=np.ceil(train_  
generator.n//train_generator.batch_size),  
                      epochs=EPOCHS)  
  
# Train the loaded model for more epochs  
history = model.fit_generator(generator=train_generator,  
                               steps_per_epoch=STEP_  
SIZE_TRAIN,  
                               epochs=EPOCHS,  
                               callbacks=[batch_gc_c  
allback, lr_finder])
```

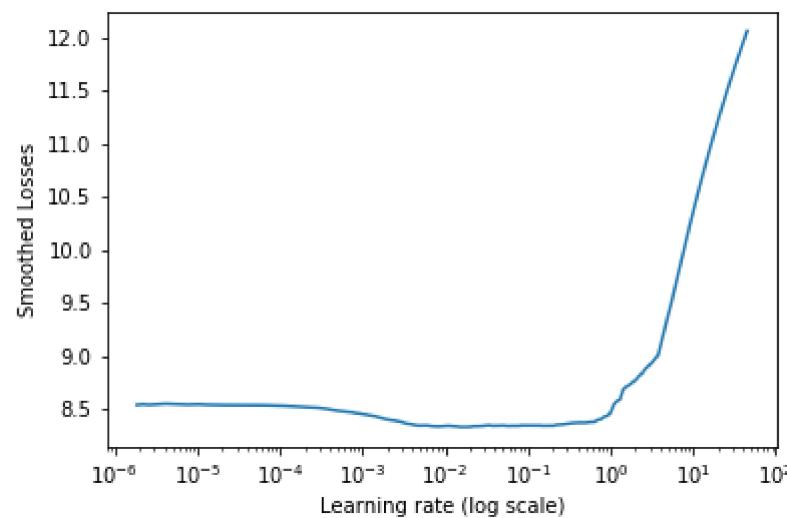
```
Epoch 1/5  
Epoch start learning rate: 1e-06  
156/156 [=====] - 217s 1s/  
step - loss: 8.4317 - acc: 0.0040 - mean_absolute_e  
rror: 3.9948e-04 - top_5_accuracy: 0.0104  
Epoch ended learning rate: 1.0  
Epoch 2/5  
Epoch start learning rate: 1.0  
49/156 [=====>.....] - ETA: 2:1  
3 - loss: 14.2739 - acc: 0.0012 - mean_absolute_err  
or: 3.9945e-04 - top_5_accuracy: 0.5459 Epoch ended  
learning rate: 70.17038
```

```
In [18]:  
lr_finder.plot_loss()
```



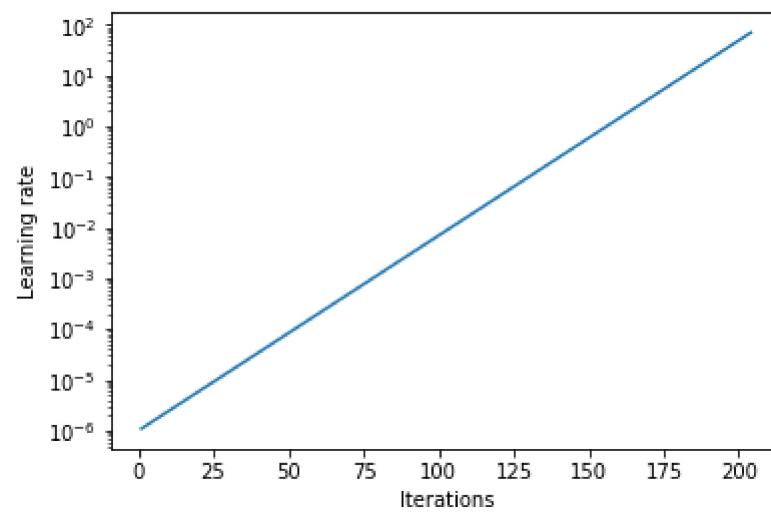
In [19]:

```
lr_finder.plot_smoothed_loss()
```



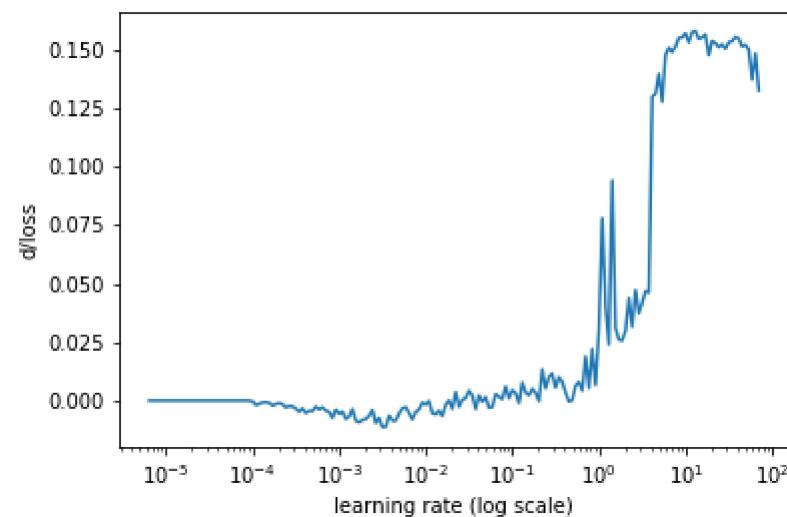
In [20]:

```
lr_finder.plot_lr()
```



In [21]:

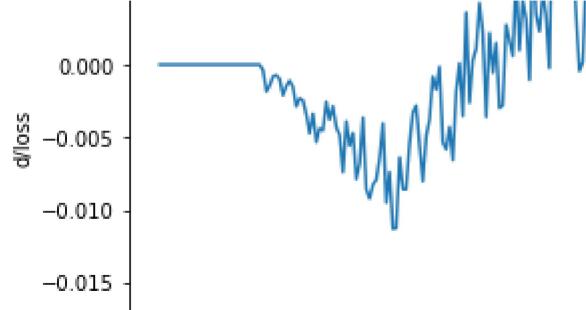
```
lr_finder.plot_loss_change(sma=50)
```



In [22]:

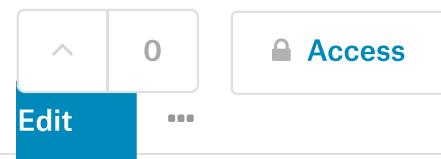
```
lr_finder.plot_loss_change(sma=50, y_lim=(-0.02, 0.01), should_lim_y=True)
```





## lr finder

Python notebook using data from [multiple data sources](#) · 23 views · multiple data sources  
 Edit tags



### Version 7

9 commits

forked from Xception

#### Notebook

#### Data

#### Output

#### Log

#### Comments

## Plot train results

```
In [23]:  
def plot_accuracy(history, should_plot_val = False,  
                  should_plot_top5 = False):  
    acc = history.history['acc']  
    l1 = plt.plot(acc, label='acc')  
  
    if should_plot_val:  
        val_acc = history.history['val_acc']  
        l2 = plt.plot(val_acc, label='val_acc')  
  
    if should_plot_top5:  
        top5_acc = history.history['top_5_accuracy']  
    ]  
    l3 = plt.plot(top5_acc, label='top_5_accuracy')  
  
    plt.legend(loc=2, fontsize="small")  
    plt.title('Model accuracy')  
    plt.ylabel('Accuracy')  
    plt.xlabel('Epoch')
```

This kernel has been released under the [Apache 2.0](#) open source license.

Did you find this Kernel useful?  
Show your appreciation with an upvote



#### Data

##### Data Sources

- ▼ Humpback Wh...
  - ▀ 7960 x 2
  - ▀ 25.4k x 2
- ▼ test.zip
  - ▀ 0027089a...
  - ▀ 00313e2d...
  - ▀ 004344e9...



### Humpback Whale Identification

Can you identify a whale by its tail?

Last Updated: 2 months ago

#### About this Competition

This training data contains thousands of images of humpback whale flukes. Individual whales have been identified by researchers and given an Id. The

008a4bc8...  
 00ac0fc...  
 00ff45291...  
 012dbdb5...

challenge is to predict the whale `Id` of images in the test set. What makes this such a challenge is that there are only a few examples for each of 3,000+ whale Ids.

Notebook

Data

Output

Log

Comments

01b1ecf7b....  
... 1000+ more  
▼ train.zip  
 002b4615...  
 00600ce1...  
 00d64188...  
 00eaedfab...  
 00fee397...  
 010a1f0eb...  
 01237f1ce...  
 01dc...  
 0202dfb2...  
 020ab0f9...  
... 1000+ more

- **train.zip** - a folder containing the training images
- **train.csv** - maps the training `Image` to the appropriate whale `Id`. Whales that are not predicted to have a label identified in the training data should be labeled as `new_whale`.
- **test.zip** - a folder containing the test images to predict the whale `Id`
- **sample\_submission.csv** - a sample submission file in the correct format

Output Files

New Dataset

New Kernel

Download All

X

#### Output Files

Model\_Xception...  
 Weights\_Xcepti...

#### About this file

This file was created from a Kernel, it does not have a description.

Model\_Xception.h5

We don't support previews for this file yet

#### Run Info

Succeeded	True	Run Time	341.3 seconds
Exit Code	0	Queue Time	0 seconds
Docker Image Name	/python(Dockerfile)	Output Size	0
Timeout Exceeded	False	Used All Space	False
Failure Message			

Log

Download Log

```
Time  Line # Log Message
2.9s    1 [NbConvertApp] Converting notebook
          __notebook__.ipynb to notebook
2.9s    2 [NbConvertApp] Executing notebook with kernel:
          python3
7.9s    3 2019-02-01 13:33:36.358697: I
          tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:964]
          successful NUMA node read from SysFS had negative value
          (-1), but there must be at least one NUMA node, so
          returning NUMA node zero
7.9s    4 2019-02-01 13:33:36.359533: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:1432]
          Found device 0 with properties:
          name: Tesla K80 major: 3 minor: 7
          memoryClockRate(GHz): 0.8235
          pciBusID: 0000:00:04.0
          totalMemory: 11.17GiB freeMemory: 11.10GiB
          2019-02-01 13:33:36.359619: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:1511]
          Adding visible gpu devices: 0
8.4s    5 2019-02-01 13:33:36.890580: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:982]
          Device interconnect StreamExecutor with strength 1
          edge matrix:
          2019-02-01 13:33:36.890652: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:988]
          0
          2019-02-01 13:33:36.890669: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:1001]
          0: N
8.4s    6 2019-02-01 13:33:36.894585: I
          tensorflow/core/common_runtime/gpu/gpu_device.cc:1115]
          Created TensorFlow device
          (/job:localhost/replica:0/task:0/device:GPU:0 with
          10758 MB memory) -> physical GPU (device: 0, name:
          Tesla K80, pci bus id: 0000:00:04.0, compute
          capability: 3.7)
338.0s   7 [NbConvertApp] Writing 503091 bytes to
          __notebook__.ipynb
340.3s   8 [NbConvertApp] Converting notebook
          __notebook__.ipynb to html
340.9s   9 [NbConvertApp] Support files will be in
          __results__files/
          [NbConvertApp] Making directory __results__files
340.9s  10 [NbConvertApp] Making directory __results__files
          [NbConvertApp] Writing 356565 bytes to
          __results__.html
340.9s  11
340.9s  13 Complete. Exited with code 0.
```

## Comments (0)



Click here to enter a comment...